

Assignment 2 Solutions

CS283, Computer Vision

1. (a) We denote

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i, \bar{y} = \frac{1}{N} \sum_{i=1}^N y_i. \quad (1)$$

It is easy to see that $x' = sx + t_x$ and $y' = sy + t_y$. From the requirement about the centroid, and using (1), we get

$$\sum_{i=1}^N x'_i = 0 \Rightarrow \sum_{i=1}^N (sx_i + t_x) = 0 \Rightarrow s \sum_{i=1}^N x_i + Nt_x = 0 \Rightarrow t_x = -s\bar{x}, \quad (2)$$

and similarly

$$t_y = -s\bar{y}. \quad (3)$$

From the average distance constraint, and using (2) and (3), we get

$$\begin{aligned} \frac{1}{N} \sum_{i=1}^N \sqrt{(sx_i + t_x)^2 + (sy_i + t_y)^2} &= \sqrt{2} \Rightarrow \\ s &= \frac{\sqrt{2}N}{\sum_{i=1}^N \sqrt{(x_i - \bar{x})^2 + (y_i - \bar{y})^2}}. \end{aligned} \quad (4)$$

Replacing (4) into (2) and (3), we get the complete expressions for t_x and t_y .

- (b) A possible implementation for this part is included below.

```
1 %%
2 % CS283 Fall 2015
3 % Student: Ioannis Gkioulekas
4 % Homework assignment 2
5 % Question 1, part b
6 %
7 % Center and scale samples
8 %
9 function T = getT(X)
10
11 if (~ismatrix(X)) || (size(X, 2) ≠ 2) || (size(X, 1) == 0)
12     error('Input matrix must be a Nx2 matrix.');
```

```

33 ty = - s * Ybar;
34
35 %% Form and return T
36 T = [s 0 tx;
37      0 s ty;
38      0 0 1];

```

(c) The resulting plot is shown in figure 1.




Figure 1: Resulting plots for script of question 3c.

Line 2 of the test script converts the inhomogeneous co-ordinates in the matrix X to their homogeneous versions by concatenating a 1 to each vector, and then applies the similarity matrix T to them. Line 3 then converts the homogeneous co-ordinates back to inhomogeneous ones, by dividing the first two coordinates with the third one.

2. (a) Each corresponding pair of points $x = [x_1, x_2, x_3]^T, y = [y_1, y_2, y_3]^T \in \mathbb{R}^3$, in homogeneous coordinates, is related by the homography H as $Hx \times y = 0$. If we define $H^T = [h_1, h_2, h_3]$ where $h_1, h_2, h_3 \in \mathbb{R}^3$ correspond to the rows of H , we can write

$$Hx = \begin{bmatrix} h_1^T x \\ h_2^T x \\ h_3^T x \end{bmatrix}, \quad (5)$$

and using this we can express the homography relationship as

$$Hx \times y = \begin{bmatrix} (h_2^T x)y_3 - (h_3^T x)y_2 \\ (h_3^T x)y_1 - (h_1^T x)y_3 \\ (h_1^T x)y_2 - (h_2^T x)y_1 \end{bmatrix} = 0. \quad (6)$$

From every corresponding pair of points, using (6) we get three equations for the elements of H , only two of which are linearly independent. If we select the first two equations, then they can equivalently be written in the form

$$\underbrace{\begin{bmatrix} 0 & 0 & 0 & y_3 x^T & -y_2 x^T \\ y_2 x^T & -y_1 x^T & 0 & 0 & 0 \end{bmatrix}}_A \begin{bmatrix} h_1 \\ h_2 \\ h_3 \end{bmatrix} = 0. \quad (7)$$

If we have multiple corresponding pairs of points, then for each of those we can form a matrix A

of the form of (7). If we stack up the matrices A for all pairs, we get

$$\underbrace{\begin{bmatrix} A_1 \\ A_2 \\ \vdots \\ A_N \end{bmatrix}}_{\Lambda} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \end{bmatrix} = 0, \quad (8)$$

where Λ is a $2N \times 9$ matrix. Adding the constraint $\sum_{i,j} H_{ij}^2 = 1$ to avoid the trivial solution $H = 0$, we can obtain the least-squares solution of the above overcomplete linear system. To do so, we can use the Singular Value Decomposition (SVD) of Λ , as was discussed in class (see also section 3.1 of [1]).

As suggested in [2], we should firstly normalize the points x and y , and then calculate the homography between their respective images. If T_x and T_y are the similarity matrices for the two sets of points, and H' the homography that sends $T_x x$ to $T_y y$, then overall transformation is given by $H = T_y^{-1} H' T_x$. As usual however, the inverse is used *only for notational purposes*, in order to express the solution in closed-form. It is almost never necessary to compute the inverse of a matrix. In practice, we can solve the above problem by using some appropriate factorization method, such as QR factorization. These methods offer significant advantages both in terms of computational efficiency and numerical stability. In Matlab, these can be easily employed using the backslash operator, as shown in the implementation in the following.

As discussed in the hints, in Matlab it is important to avoid loops whenever possible. In the implementation below, we do provide a version that uses no loops whatsoever. If you need to use a loop, there are several ways in which you can improve performance. Possibly the most important is to make sure that you have *preallocated* all memory you are going to use inside the loop. The difference in performance with and without preallocation can be dramatic. You can read more about this technique in Matlab's Techniques for Improving Performance¹, as well as in the Undocumented Matlab website². In the following implementation, we also provide a loopy version of the code, where we take care to preallocate the matrix A before entering the for loop. Below is a possible implementation of function `getH.m`, following the discussion above. It uses `getT.m` from question 4.

```

1  %%
2  % CS283 Fall 2015
3  % Student: Ioannis Gkioulekas
4  % Homework assignment 2
5  % Question 2, part a
6  %
7  % Estimate a homography from corresponding pairs of points
8  %
9  function H = getH(X1, X2)
10
11  if (~ismatrix(X1)) || (~ismatrix(X2)) || ~isempty(find(size(X1) ~= size(X2), 1))...
12      || (size(X1, 2) ~= 2) || (size(X1, 1) == 0)
13      error('Input must be two Nx2 matrices.');
```

¹http://www.mathworks.com/help/matlab/matlab_prog/techniques-for-improving-performance.html

²<http://undocumentedmatlab.com/blog/preallocation-performance/>

```

24
25 %% Construct matrix A of the homogeneous linear system
26 % 1) Loop-free version.
27 % A = [zeros(N, 3), -repmat(X2h(3, :)', [1 3]) .* X1h', repmat(X2h(2, :)', [1 ...
      3]) .* X1h'];
28 % repmat(X2h(3, :)', [1 3]) .* X1h', zeros(N, 3), -repmat(X2h(1, :)', [1 3]) ...
      .* X1h'];
29
30 % 2) Loopy version, two variants.
31 % Make sure to preallocate A.
32 A = zeros(2 * N, 9);
33 % A = zeros(3 * N, 9); % for redundant version where A is 3*N x 9 matrix.
34 for iter = 1:N,
35     A((iter * 2 - 1):(iter * 2), :) = ...
36     [zeros(1, 3) -X2h(3, iter) * transpose(X1h(:, iter)) X2h(2, iter) * ...
      transpose(X1h(:, iter));
37     X2h(3, iter) * transpose(X1h(:, iter)) zeros(1, 3) -X2h(1, iter) * ...
      transpose(X1h(:, iter))];
38 % Alternatively, for redundant version where A is 3*N x 9 matrix.
39 % A((iter * 3 - 2):(iter * 3), :) = ...
40 % [zeros(1, 3) -X2h(3, iter) * transpose(X1h(:, iter)) X2h(2, iter) * ...
      transpose(X1h(:, iter));
41 % X2h(3, iter) * transpose(X1h(:, iter)) zeros(1, 3) -X2h(1, iter) * ...
      transpose(X1h(:, iter))];
42 % -X2h(2, iter) * transpose(X1h(:, iter)) X2h(1, iter) * transpose(X1h(:, ...
      iter)) zeros(1, 3)];
43 end;
44
45 %% Compute SVD and take last vector of V
46 [U, S, V] = svd(A);
47 h = V(:, end);
48 Htild = transpose(reshape(h, [3 3]));
49
50 %% "Remove" normalizations
51 % We use the backslash operator (which employs QR decomposition) instead of
52 % calling inv or using H ^ (-1).
53 H = T2 \ Htild * T1;
54
55
56 function Hcoords = in2hom(Icoords)
57
58 Hcoords = [Icoords; ones(1, size(Icoords, 2))];

```

(b) A possible implementation of function `applyH.m` is shown below.

```

1 %%
2 % CS283 Fall 2015
3 % Student: Ioannis Gkioulekas
4 % Homework assignment 2
5 % Question 2, part b
6 %
7 % Apply a homography to an input image
8 %
9 function Iout = applyH(Iin, H)
10
11 % convert to double
12 Iin = im2double(Iin);
13
14 % find borders
15 w = size(Iin, 2);
16 h = size(Iin, 1);
17 corners = [1 1; 1 h; w 1; w h];
18 cornersH = transpose(H * transpose([corners ones(4, 1)]));
19 cornersH = cornersH(:, 1:2) ./ repmat(cornersH(:, 3), [1 2]);
20 xmin = min(cornersH(:, 1));
21 ymin = min(cornersH(:, 2));
22 xmax = max(cornersH(:, 1));

```

```

23 ymax = max(cornersH(:, 2));
24
25 % Set up resolution so that both images have approximately
26 % the same number of pixels.
27 xres = (xmax - xmin + 1);
28 yres = (ymax - ymin + 1);
29 zfactr = sqrt(w * h / xres / yres);
30 xres = ceil(xres * zfactr);
31 yres = ceil(yres * zfactr);
32
33 % Compute grid and corresponding co-ordinates
34 [x y] = meshgrid(linspace(xmin, xmax, xres), linspace(ymin, ymax, yres));
35 X2h = [x(:) y(:) ones(numel(x), 1)];
36 % We use the backslash operation (which employs QR decomposition) instead of
37 % calling inv or using H ^ (-1).
38 X1h = transpose(H \ transpose(X2h));
39 X1 = X1h(:, 1:2) ./ repmat(X1h(:, 3), [1 2]);
40
41 % Form image
42 numC = size(Iin,3);
43 Iout = zeros([size(x) numC]);
44 for i = 1:numC
45     Iout(:, :, i) = reshape(interp2(Iin(:, :, i), X1(:, 1), X1(:, 2), 'linear'), ...
46                               size(x));
47 end;
48 Iout(isnan(Iout)) = 0;

```

(c) A rectification of the Maxwell-Dworkin picture using the code above is shown in figure 2.

- Below is one possible solution for question 3. The function uses the coordinate system of the central image, and projects all other images to that through the use of homographies. When regions overlap, blending is done by averaging intensity estimates from all the images.

```

1 %%
2 % CS283 Fall 2015
3 % Student: Ioannis Gkioulekas
4 % Homework assignment 2
5 % Question 3
6 %
7 % Stitch together images given a set of correspondences
8 %
9 function Iout = stitch(cimg, imgs, Hs)
10
11 numI = length(imgs);
12 cimg = im2double(cimg);
13
14 for i = 1:numI
15     imgs{i} = im2double(imgs{i});
16 end;
17
18 % Compute full extents
19 w = size(cimg, 2);
20 h = size(cimg, 1);
21 extns = [1 1; 1 h; w 1; w h];
22
23 for i = 1:numI
24     w = size(imgs{i}, 2);
25     h = size(imgs{i}, 1);
26     crns = [1 1; 1 h; w 1; w h];
27     crnsH = transpose(Hs{i} * transpose([crns ones(4, 1)]));
28     crns = crnsH(:, 1:2) ./ repmat(crnsH(:, 3), [1 2]);
29     extns = [extns; crns];
30 end;
31
32 xmin = min(extns(:, 1));

```



Figure 2: Rectification of the Maxwell-Dworkin picture.

```

33 xmax = max(extns(:, 1));
34 ymin = min(extns(:, 2));
35 ymax = max(extns(:, 2));
36
37 % Compute grid and corresponding co-ordinates
38 [x y] = meshgrid([floor(xmin):ceil(xmax)], [floor(ymin):ceil(ymax)]);
39 X2h = [x(:) y(:) ones(numel(x),1)];
40
41 % Form Image
42 numC = size(cimg, 3);
43 Iout = zeros([size(x) numC]);
44 blendmap = Iout;
45
46 % — From central image
47 for i = 1:numC
48 Iout(:, :, i) = reshape(interp2(cimg(:, :, i), x(:), y(:)), size(x));
49 5
50 end;
51 blendmap(~isnan(Iout)) = 1;
52 Iout(isnan(Iout)) = 0;
53
54 % — From other images
55 for j = 1:numI
56 X1h = transpose(Hs{j} \ transpose(X2h));
57 X1 = X1h(:, 1:2) ./ repmat(X1h(:, 3), [1 2]);
58 temp_o = zeros([size(x) numC]);
59
60 for i = 1:numC
61 temp_o(:, :, i) = reshape(interp2(imgs{j}(:, :, i), X1(:, 1), X1(:, 2), linear), ...
size(x));
62 end;
63 blendmap(~isnan(temp_o)) = blendmap(~isnan(temp_o)) + 1;
64 temp_o(isnan(temp_o)) = 0;
65 Iout = Iout + temp_o;
66 end;
67 Iout = Iout ./ max(blendmap, 1);

```

4. An implementation of the robust version of `getH.m` that uses RANSAC is shown below.

```

1 %%
2 % CS283 Fall 2015
3 % Student: Ioannis Gkioulekas
4 % Homework assignment 2
5 % Question 4
6 %
7 % Estimate a homography from correspondences found using RANSAC
8 %
9 function H = ransacGetH(X1, X2)
10
11 % threshold for inliers
12 thresh = 3;
13 N = size(X1, 1);
14 X1h = [X1 ones(N, 1)];
15 X2h = [X2 ones(N, 1)];
16 numi_o = 1;
17 meane_o = 0;
18 ilist_o = [];
19 i = 0;
20
21 num_samples = 10 ^ 15;
22 while i < num_samples
23
24     % Estimate H from 4 random points
25     idx = randperm(N, 4);
26     H = getH(X1(idx, :), X2(idx, :));
27

```

```

28 % Compute distance, number of inliers, etc
29 xx2h = transpose(H * transpose(X1h));
30 xx2 = xx2h(:, 1:2) ./ repmat(xx2h(:, 3), [1 2]);
31 xx1h = transpose(H \ transpose(X2h));
32 xx1 = xx1h(:, 1:2) ./ repmat(xx1h(:, 3), [1 2]);
33 dists = sqrt(sum((X2 - xx2) .^ 2, 2)) + sqrt(sum((X1 - xx1) .^ 2, 2));
34 ilist = find(dists < 2 * thresh);
35 numi = length(ilist);
36 meane = mean(dists(numi));
37
38 % Update number of samples
39 i = i + 1;
40 num_samples = log(0.001) / log(1 - (numi / N) ^ 4);
41
42 % Update if answer s optimal
43 if numi < numi_o
44     continue;
45 end;
46
47 if (numi == numi_o) && (meane > meane_o)
48     continue;
49 end;
50
51 ilist_o = ilist;
52 meane_o = meane;
53 numi_o = numi;
54 end;
55
56 H = getH(X1(ilist, :), X2(ilist, :));

```

In figure 3, we show the image resulting from stitching the three images provided together, using the functions posted above.



Figure 3: Stitched images as produced by the functions for questions 3 and 4.

References

- [1] FORSYTH, D., AND PONCE, J. *Computer Vision: A modern approach*. Prentice Hall, 2003.
- [2] HARTLEY, R., AND ZISSERMAN, A. *Multiple view geometry in computer vision*. Cambridge Univ Press, 2003.